

# Mohammed Adil Siraju

[mohdadilsiraju@gmail.com](mailto:mohdadilsiraju@gmail.com) | [adilsiraju.vercel.app](https://adilsiraju.vercel.app) | [linkedin.com/in/adilsiraju](https://linkedin.com/in/adilsiraju) | [github.com/adilsiraju](https://github.com/adilsiraju) | +91-8113936995

## Machine Learning Engineer | Deep Learning | NLP | Computer Vision

Indian Citizen | Open to Relocation (UAE / India tech hubs)



### Professional Summary

Machine Learning & Deep Learning Engineer with a B.E. in AIML (First Class with Distinction). Skilled in PyTorch, FastAI, Hugging Face, and building end-to-end ML pipelines including data preprocessing, model training, hyperparameter tuning, evaluation, and deployment. Hands-on experience in computer vision, NLP, and MLOps tools. Strong track record of delivering deployed ML apps with real-world usage.

### Education

Bachelor of Engineering – Artificial Intelligence & Machine Learning (FCD, CGPA: 7.33/10)

Dec 2021 – Aug 2025

P A College of Engineering,

Mangalore

Relevant Coursework: Artificial Intelligence, Machine Learning, Data Structures, Deep Learning, Cloud Computing, Data Science, Algorithms, Operating Systems

### Skills

<b>Core ML/DL:</b>	Computer Vision, NLP, Supervised & Unsupervised Learning, Transfer Learning, Model Deployment
<b>Frameworks/Libraries:</b>	PyTorch, FastAI, Hugging Face Transformers, scikit-learn
<b>MLOps &amp; Tools:</b>	Docker, Git, Kubernetes (basic), Gradio, Jupyter
<b>Programming:</b>	Python, SQL

### Experience

**AI DevOps Engineer** - Intern, Rooman Technology – Mangalore

Sep 2024 – Feb 2025

- Set up CI/CD pipelines in lab environments to automate ML model testing & deployment.
- Worked with Docker and Kubernetes for containerized model deployment simulations.
- Supported team in designing scalable infrastructure for ML workflows.

### Projects

**Medical Document Classifier** | *Hugging Face Transformers, Gradio, Python*

[Medical-Case-Classifier](#)

- Developed a medical document classifier by fine-tuning a pre-trained **Bio\_ClinicalBERT** model on a dataset of clinical texts.
- Achieved** a test accuracy of **81.08%** in classifying documents across **13+ medical specialties**, with a top F1-score of **0.968** for the "Cardiovascular / Pulmonary" specialty.
- Deployed an interactive web application using Gradio, demonstrating an end-to-end solution for real-time medical text classification.

**Architectural Style Classifier** | *FastAI, PyTorch, Gradio*

[Architecture-Style-Classifer](#)

- Built and deployed a CNN-based classifier to identify 25 architectural styles with 73% accuracy.
- Trained on 10k+ images from Kaggle across 25 architectural styles, with custom dataset curation and preprocessing.
- Fine-tuned a pretrained model via Transfer Learning and deployed via Gradio web app (with Gradio API) and GitHub Pages site, enabling real-time inference for 500+ public users.

**Netflix Userbase Analysis** | *Python, Pandas, Matplotlib, Seaborn*

[Netflix-Userbase-EDA](#)

- Conducted EDA on Netflix userbase (8k+ records) to uncover insights on demographics, subscriptions, and revenue distribution.
- Built visualizations using Pandas/Seaborn to analyze age, gender, and country-level revenue patterns.
- Identified churn-prone user groups and regional trends, providing business implications for marketing optimization and revenue growth.

### Certifications

- Python for Data Science | IBM (2022)
- Problem Solving Using Computational Thinking | University of Michigan (2022).
- CS50x - Computer Science Fundamentals | Harvard University (2025)